


The Pipeline




- We have looked at importing, exploring and transforming data.
- Now it is time to do something with our data: **MODELS!**
- Here we discuss one of the most straight forward machine learning models: **the decision tree.**

Machine Learning

- What is Machine Learning?
 - Programs that get better with *experience* given a *task* and some *performance measure*.
 - Learning to classify news articles
 - Learning to recognize spoken words
 - Learning to play board games
 - Learning to navigate (e.g. self-driving cars)
- Usually involves some sort of inductive reasoning step.

Inductive Reasoning

- Deductive reasoning (rule based reasoning)
 - From the general to the specific
- Inductive reasoning
 - From the specific to the general



Note: not to be confused with mathematical induction!

Example - Deduction

- Rules:
 - If Betty wears a white dress then it is Sunday.
 - Betty wears a white dress.
- Deductive step:
 - You infer or *deduce* that today is Sunday.

If Betty wears a white dress then it is Sunday.
Betty wears a white dress.



Today is Sunday.

Deduction

Inference is the act or process of drawing a conclusion based solely on what one already knows.

Example - Induction

- Facts: every time you see a swan you notice that the swan is white.
- Inductive step: you infer that all swans are white.

All Swans
are white.



Induction


Observed Swans
are white.

Inference is the act or process of drawing a conclusion based solely on what one already knows.

Observation

- Deduction is “truth preserving”
 - If the rules employed in the deductive reasoning process are sound, then, what holds in the theory will hold for the deduced facts.
- Induction is NOT “truth preserving”
 - It is more of a statistical argument
 - The more swans you see that are white, the more probable it is that all swans are white..
But this does not exclude the existence of black swans

Observation



$D \equiv$ observations

$X \equiv$ universe of all swans

Different Styles of Machine Learning

- Supervised Learning
 - The learning needs explicit examples of the concept to be learned (e.g. white swans, playing tennis, *etc*)
- Unsupervised Learning
 - The learner discovers autonomously any structure in a domain that might represent an interesting concept

Knowledge - Representing what has been learned

- Symbolic Learners (transparent models)
 - If-then-else rules
 - Decision trees
 - Association rules
- Sub-Symbolic Learners (non-transparent models)
 - (Deep) Neural Networks
 - Clustering (Self-Organizing Maps, k-Means)
 - Support Vector Machines

Decision Trees

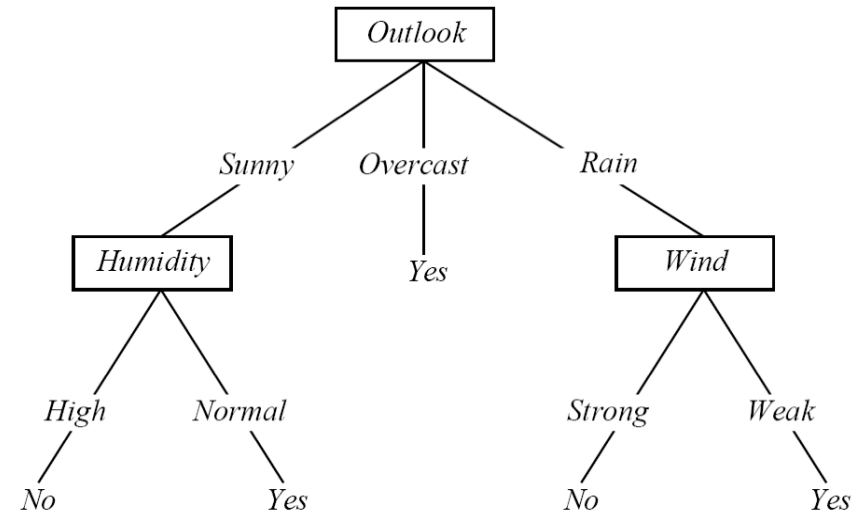
- Learn from labeled observations - supervised learning
- Represent the knowledge learned in form of a tree

Example: learning when to play tennis.

- Examples/observations are days with their observed characteristics and whether we played tennis or not

Play Tennis Example


| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |




Decision Tree Learning

| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Cool | Normal | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |

Facts or Observations



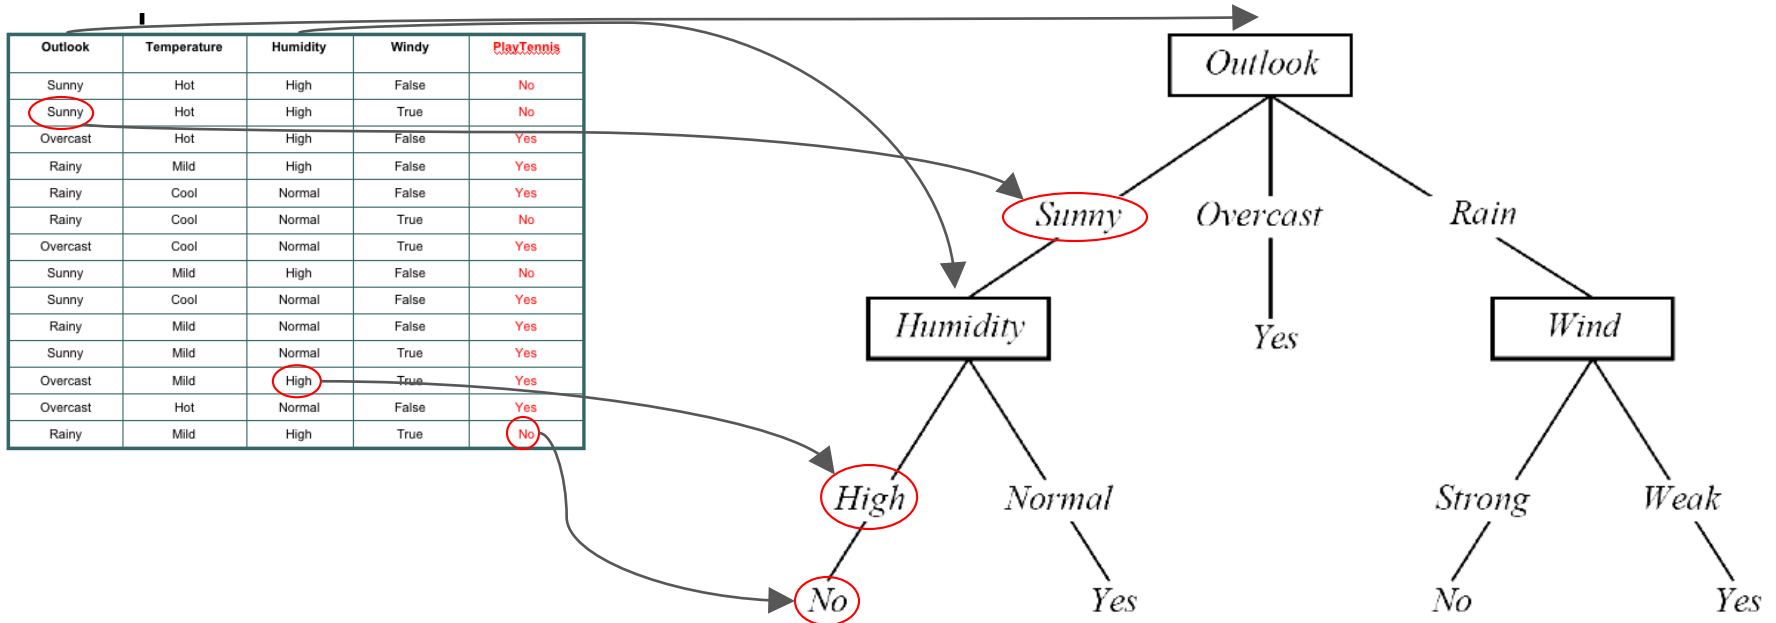
Induction



Theory

Interpreting a DT

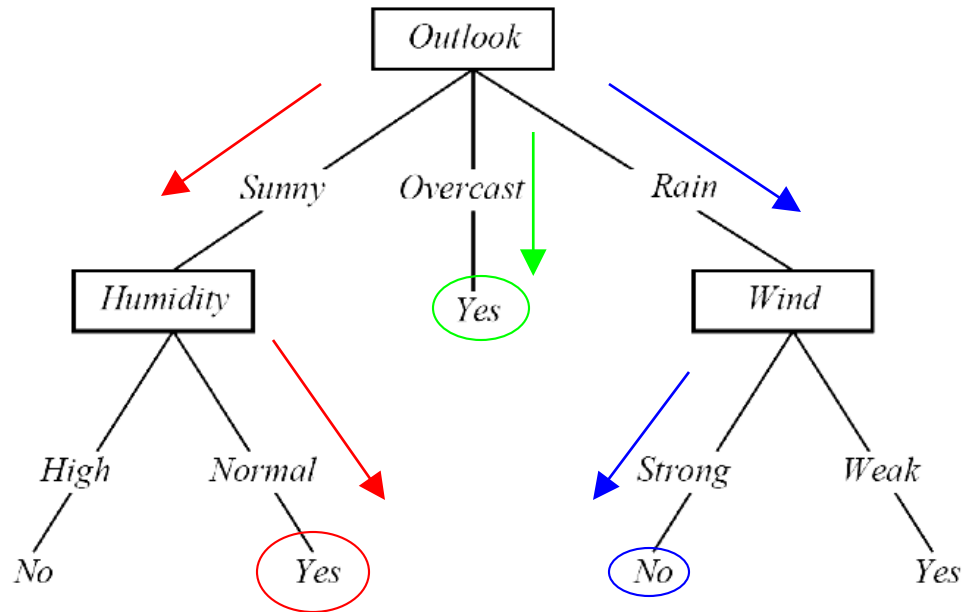
DT \equiv Decision
Tree



- A DT uses the features of an observation table as nodes and the feature values as links.
- All feature values of a particular feature need to be represented as links.
- The target feature is special - its values show up as leaf nodes in the DT.

Interpreting a DT

Each path from the root of the DT to a leaf can be interpreted as a decision rule.



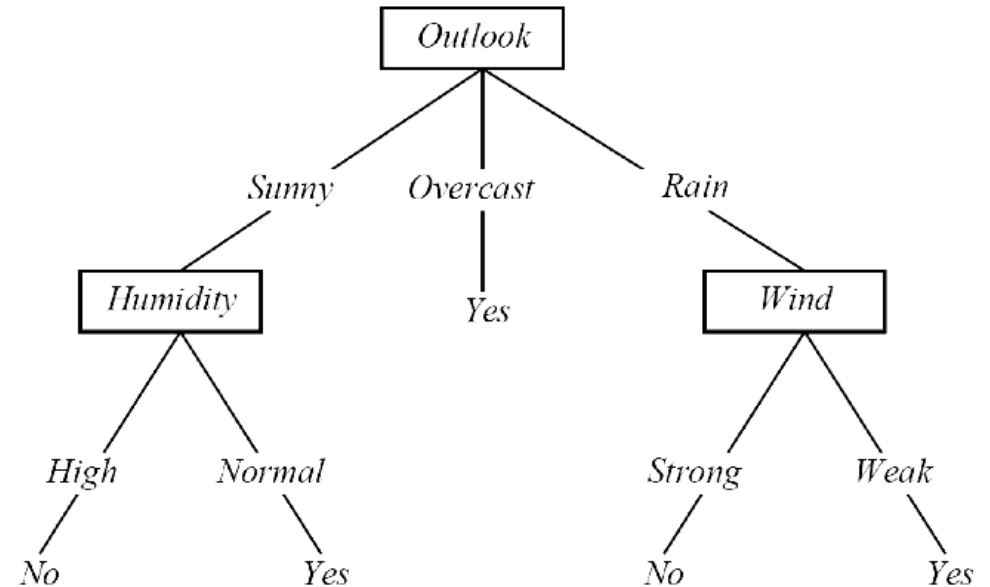
IF *Outlook = Sunny* AND *Humidity = Normal* THEN *Playtennis = Yes*

IF *Outlook = Overcast* THEN *Playtennis = Yes*

IF *Outlook = Rain* AND *Wind = Strong* THEN *Playtennis = No*

DT: Explanation & Prediction

| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |



Explanation: the DT summarizes (explains) all the observations in the table perfectly \Rightarrow 100% Accuracy

Prediction: once we have a DT (or model) we can use it to make predictions on observations that are not in the original training table, consider:

Outlook = Sunny, Temperature = Mild, Humidity = Normal, Windy = False, Playtennis = ?

Constructing DTs

- How do we choose the attributes and the order in which they appear in a DT?
 - Recursive partitioning of the original data table
 - Heuristic - each generated partition has to be “less random” (entropy reduction) than previously generated partitions


Entropy

- S is a sample of training examples
- p^+ is the proportion of positive examples in S
- p^- is the proportion of negative examples in S
- Entropy measures the impurity (randomness) of S

S {

| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |

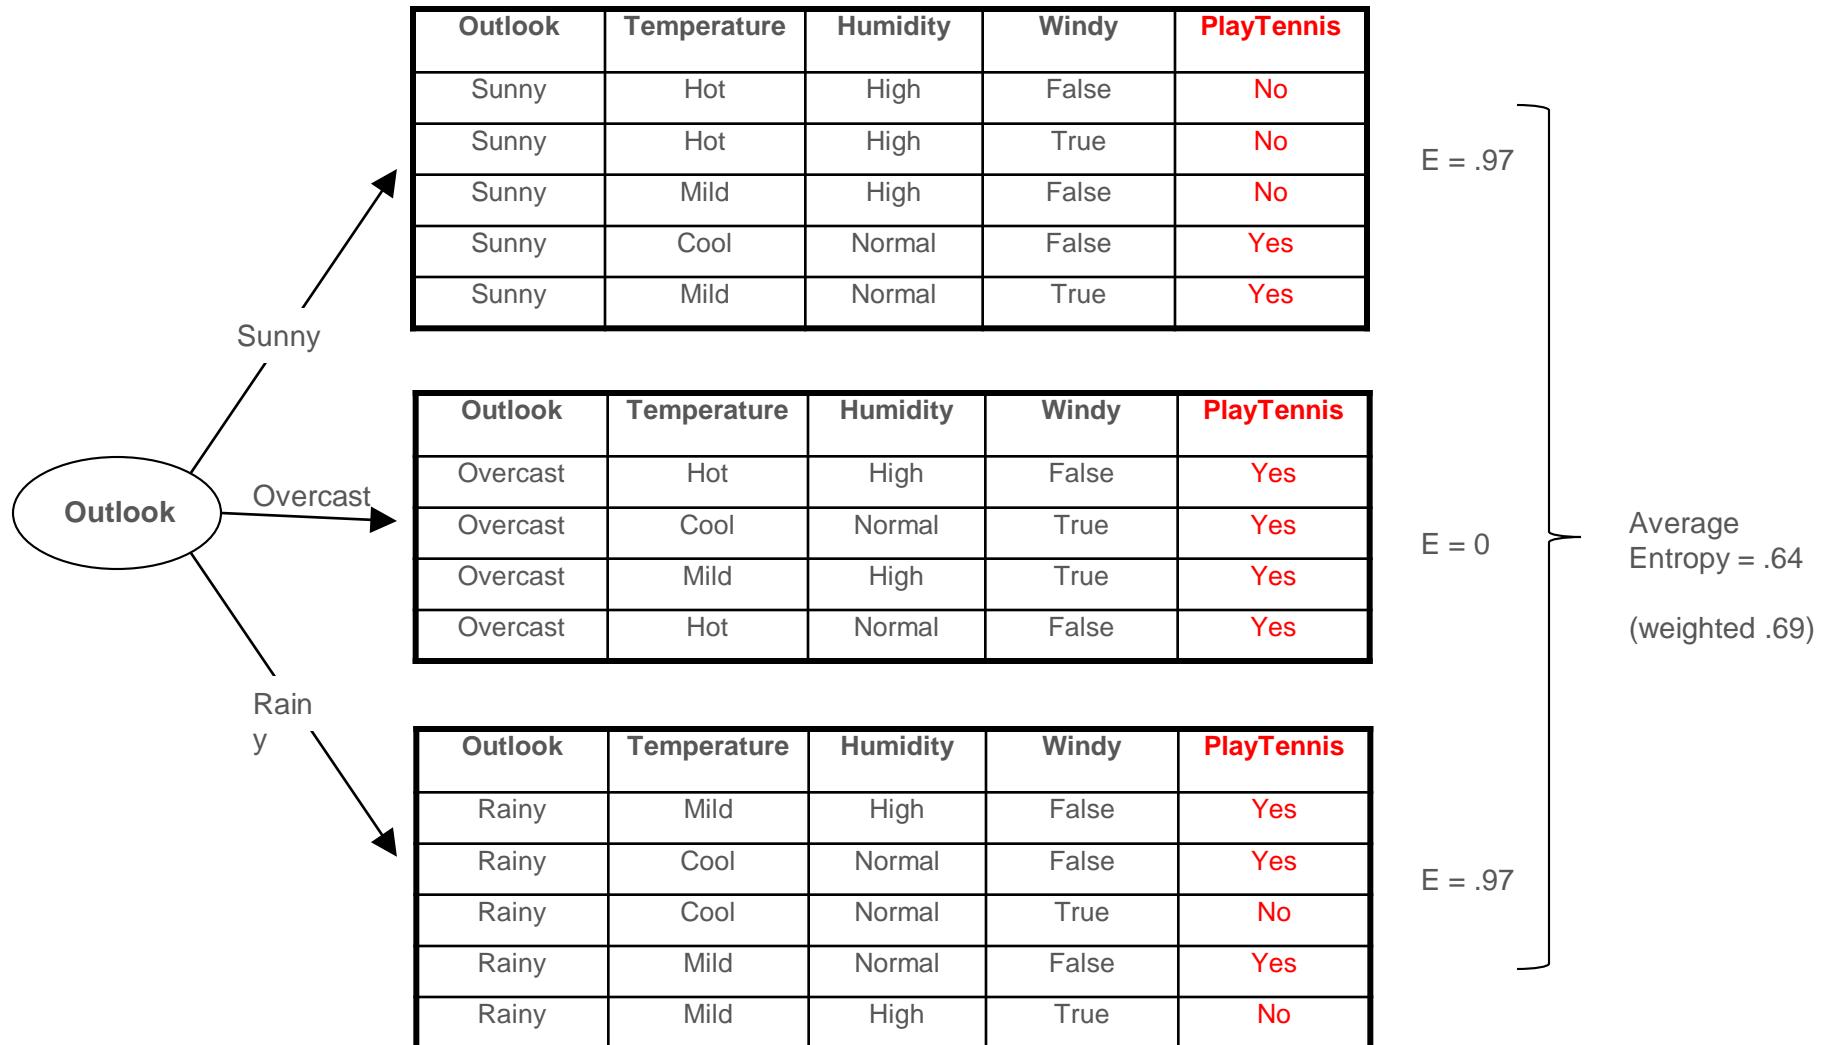
$$Entropy(S) = Entropy([9+,5-]) = .94$$



$$Entropy(S) \equiv - p^+ \log_2 p^+ - p^- \log_2 p^-$$

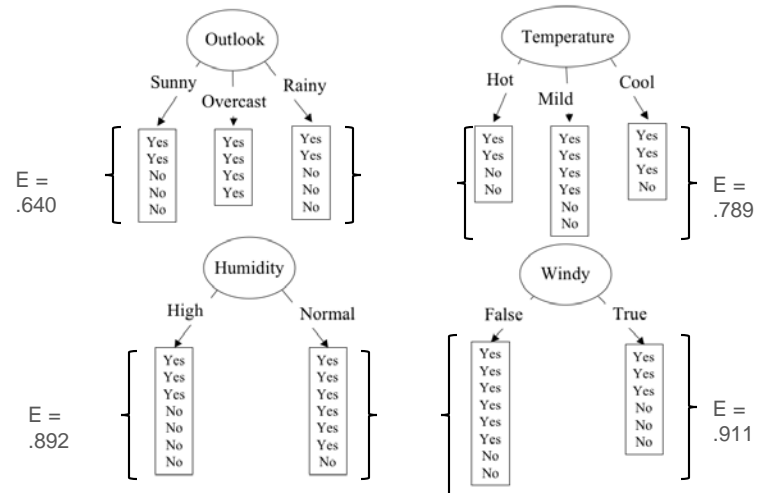
$$\text{AvgEntropy}(S, A) = \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} E(S_v) \quad (\text{weighted average})$$

Partitioning the Data Set



Partitioning in Action

| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |



Recursive Partitioning

Partition(*Examples*, *TargetAttribute*, *Attributes*)

Examples are the training examples. *TargetAttribute* is a binary (+/-) categorical dependent variable and *Attributes* is the list of independent variables which are available for testing at this point. This function returns a decision tree.

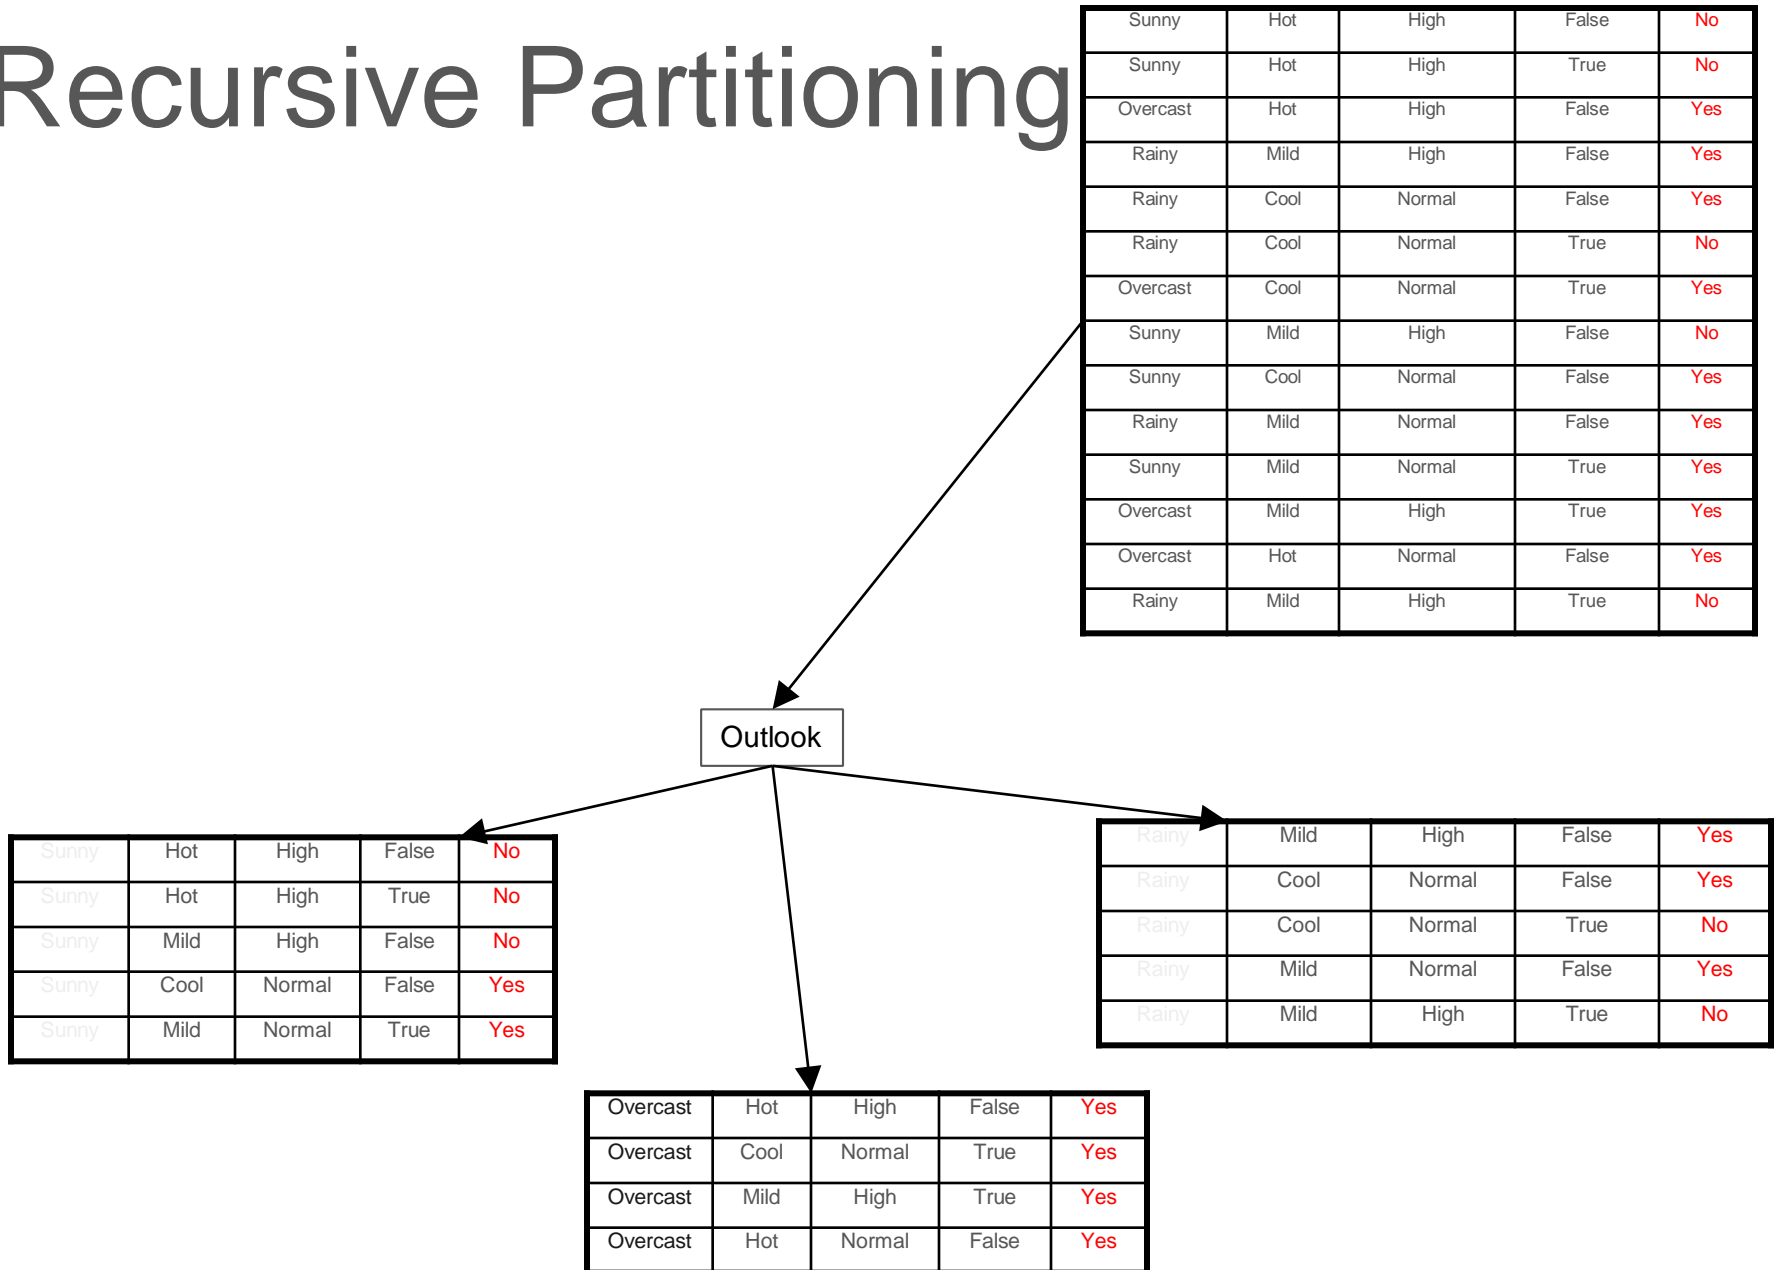
- Create a *Root* node for the tree.
- If all *Examples* are positive then return *Root* as a leaf node with label = +.
- Else if all *Examples* are negative then return *Root* as a leaf node with label = -.
- Else if *Attributes* is empty then return *Root* as a leaf node with label = most common value of *TargetAttribute* in *Examples*.
- Otherwise
 - $A :=$ the attribute from *Attributes* that reduces entropy the most on the *Examples*.
 - $Root := A$
 - For each $v \in \text{values}(A)$
 - Add a new branch below the *Root* node with value $A = v$
 - Let $Examples_v$ be the subset of *Examples* where $A = v$
 - If $Examples_v$ is empty then add new leaf node to branch with label = most common value of *TargetAttribute* in *Examples*.
 - Else add new subtree to branch
 $\text{Partition}(Examples_v, \text{TargetAttribute}, \text{Attributes} - \{A\})$
- Return *Root*

Recursive Partitioning

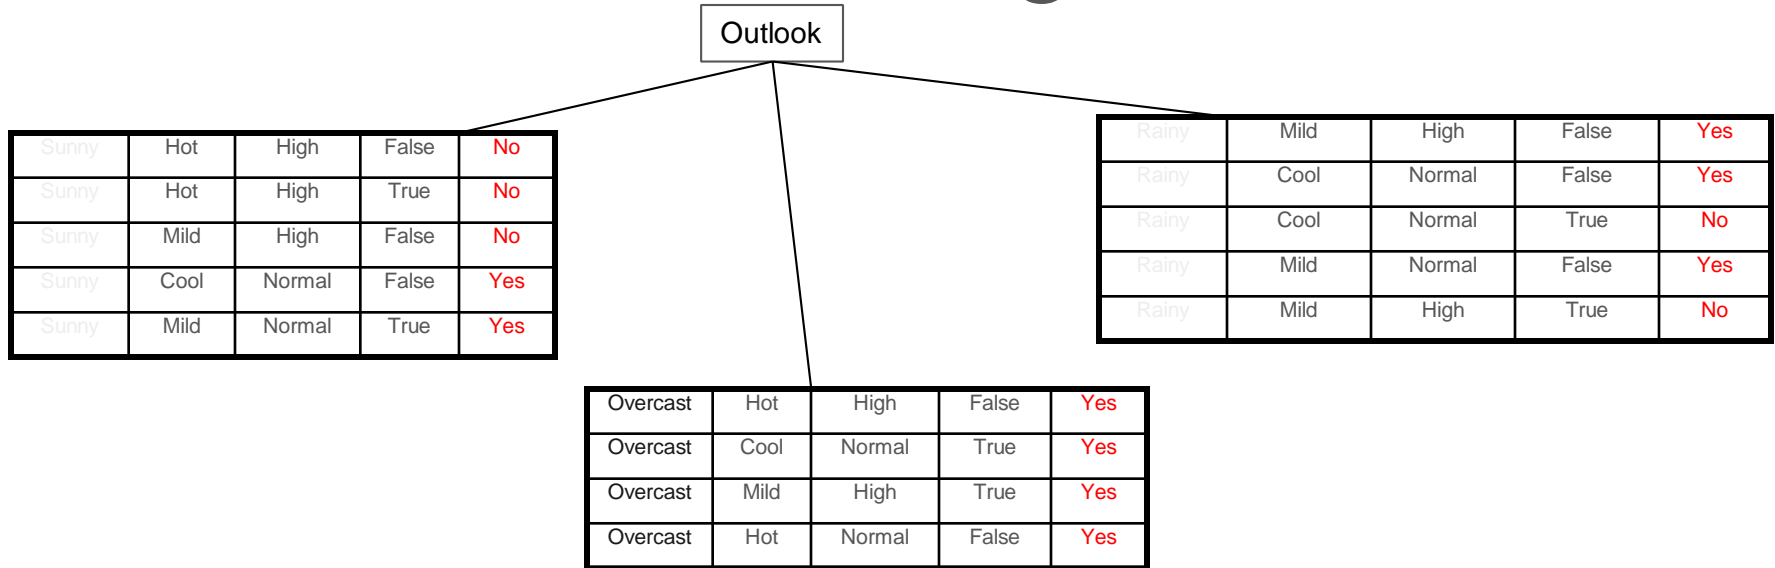
Our data set:

| Outlook | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Overcast | Hot | High | False | Yes |
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Overcast | Cool | Normal | True | Yes |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |
| Rainy | Mild | High | True | No |

Recursive Partitioning



Recursive Partitioning



Recursive Partitioning

Outlook

| | | | | |
|-------|------|--------|-------|-----|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |

Humidity

| | | | | |
|----------|------|--------|-------|-----|
| Overcast | Hot | High | False | Yes |
| Overcast | Cool | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |

| | | | | |
|-------|------|--------|-------|-----|
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Rainy | Mild | Normal | False | Yes |
| Rainy | Mild | High | True | No |

| | | | | |
|-------|------|--------|-------|-----|
| Sunny | Cool | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |

| | | | | |
|-------|------|------|-------|----|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Sunny | Mild | High | False | No |

Recursive Partitioning

Outlook

| | | | | |
|-------|------|--------|-------|-----|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Sunny | Mild | High | False | No |
| Sunny | Cool | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |

Humidity

| | | | | |
|----------|------|--------|-------|-----|
| Overcast | Hot | High | False | Yes |
| Overcast | Cool | Normal | True | Yes |
| Overcast | Mild | High | True | Yes |
| Overcast | Hot | Normal | False | Yes |

| | | | | |
|-------|------|--------|-------|-----|
| Sunny | Cool | Normal | False | Yes |
| Sunny | Mild | Normal | True | Yes |

| | | | | |
|-------|------|--------|-------|-----|
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Cool | Normal | True | No |
| Rainy | Mild | Normal | False | Yes |
| Rainy | Mild | High | True | No |

Windy

| | | | | |
|-------|------|--------|-------|-----|
| Rainy | Mild | High | False | Yes |
| Rainy | Cool | Normal | False | Yes |
| Rainy | Mild | Normal | False | Yes |

| | | | | |
|-------|------|--------|------|----|
| Rainy | Cool | Normal | True | No |
| Rainy | Mild | High | True | No |

| | | | | |
|-------|------|------|-------|----|
| Sunny | Hot | High | False | No |
| Sunny | Hot | High | True | No |
| Sunny | Mild | High | False | No |

Machine Learning in Python - Scikit-Learn

- We will be using the Scikit-Learn module to build decision trees.
 - Scikit-learn or sklearn for short provides all kinds of models
 - Neural networks
 - Support vector machines
 - Clustering algorithms
 - Linear regression
 - *etc*
- We will be using the treeviz module to visualize decision trees.
 - A simple ASCII based tree visualizer

SKlearn Decision Tree Basics

Training data needs to be structured into a *feature matrix* and a *target vector*.

In the feature matrix one row for each observations.

In the target vector one entry for each observation.

NOTE: rows and vector entries have to be consistent!

